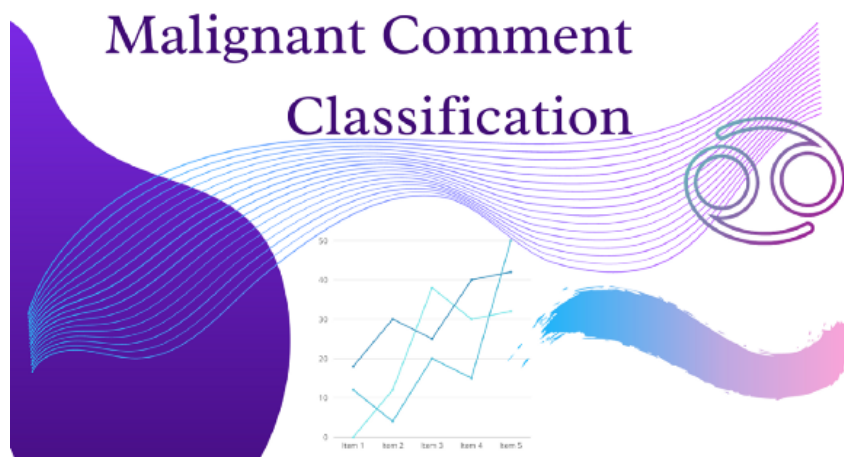


MALIGNANT COMMENTS CLASSIFICATION PROJECT



Online forums and social media platforms have provided individuals with the means to put forward their thoughts and freely express their opinion on various issues and incidents. In some cases, these online comments contain explicit language which may hurt the readers. Comments containing explicit language can be classified into myriad categories such as Toxic, Severe Toxic, Obscene, Threat, Insult, and Identity Hate. The threat of abuse and harassment means that many people stop expressing themselves and give up on seeking different opinions.

Step 1: Importing Libraries :

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

import nltk
import string
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
import nltk
from nltk.stem import PorterStemmer, WordNetLemmatizer
```

```

from wordcloud import WordCloud
from sklearn.feature_extraction.text import TfidfVectorizer
import xgboost
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score,
confusion_matrix,
classification_report, roc_curve, roc_auc_score, auc
from sklearn.model_selection import train_test_split
from sklearn.metrics import
accuracy_score, classification_report, confusion_matrix, f1_score
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import
cross_val_score, GridSearchCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import
RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier

import joblib

import warnings
warnings.filterwarnings('ignore')

```

Step 2: Importing Dataset :

```

df_train =
pd.read_csv("/content/drive/MyDrive/train.csv", encoding='utf-8')

df_test =
pd.read_csv("/content/drive/MyDrive/test.csv", encoding='utf-8')

df_train.head()

```

	id	comment_text	malignant	highly_malignant	rude	threat	abuse	loathe
0	0000997932d777bf	Explanation\n\nWhy the edits made under my usern...	0	0	0	0	0	0
1	000103f0d9c9b60f	D'awwl He matches this background colour I'm s...	0	0	0	0	0	0
2	000113f07ec002fd	Hey man, I'm really not trying to edit war. IT...	0	0	0	0	0	0
3	0001b41b1c6bb37e	"\nMore\nI can't make any real suggestions on ...	0	0	0	0	0	0
4	0001d958c54c5e35	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0

```

print(df_train.shape)
print(df_test.shape)

(159571, 8)
(153164, 2)

print(df_train.info())
print(df_test.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 159571 entries, 0 to 159570
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                    159571 non-null object
1   comment_text          159571 non-null object
2   malignant             159571 non-null int64
3   highly_malignant      159571 non-null int64
4   rude                 159571 non-null int64
5   threat               159571 non-null int64
6   abuse                159571 non-null int64
7   loathe               159571 non-null int64
dtypes: int64(6), object(2)
memory usage: 9.7+ MB
None
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 153164 entries, 0 to 153163
Data columns (total 2 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                    153164 non-null object
1   comment_text          153164 non-null object
dtypes: object(2)
memory usage: 2.3+ MB
None

```

Step 3: Problem Statement and Background

The background for the problem originates from the multitude of online forums, where-in people participate actively and make comments. As the comments some times may be abusive, insulting or even hate-based, it becomes the responsibility of the hosting organizations to ensure that these conversations are not of negative type. The task was thus to build a model which could make prediction to classify the comments into various categories. Through the `head().method` we can observed 8 rows. As per problem statement this is classification problem and here need to predict whether comment is malignant or not.







Step 4: Studying data & identifying hidden patterns

I had a dataset of 159571 samples of comments along with their labels. I observed that every 1 in 10 samples was toxic, every 1 in 50 samples was obscene and insulting, but the occurrences of sample being severe-toxic, threat and identity hate was extremely rare. The first 5 rows appeared as shown above.


For the next visualisation, We have checked the null values also and I had length of comments on the independent axis again similar to the previous plot. But instead of counting number of comments, I counted comments belonging to each of the different categories.

```
#checking the values counts
columns = [ 'malignant', 'highly_malignant', 'rude',
            'threat',
            'abuse', 'loathe']

for col in columns:
    print(col)
    print(df_train[col].value_counts())
    sns.countplot(df_train[col])
    plt.show()
```



```
# checking the most offensive words
hams = df_train['comment_text'][df_train['malignant']==1]
spam_cloud = WordCloud(width=600, height=400,
background_color = 'black', max_words=50).generate(''.
join(hams))
plt.figure(figsize=(10,8), facecolor='k')
plt.imshow(spam_cloud)
plt.axis('off')
plt.tight_layout(pad=0)
plt.show()
```



```
# checking the total distribution of comments
df_distribution = df_train[columns].sum()\
                  .to_frame()\
                  .rename(columns={0: 'count'})\
                  .sort_values('count')
df_distribution.plot.pie(y = 'count',
                        title = 'Label distribution over
comments',
                        figsize=(5,5))\
                        .legend(loc = 'center left', bbox_to_anchor = (1.3, 0.5))
```

STEP 5: Data Preprocessing

1. removal of special character marks: There are many numeric and special characters which we have to remove from dataset.

```
#replacing email address with email

df_train['comment_text']
=df_train['comment_text'].str.replace(r'^.+@[^\.]*. *[a-z]
{2,}$', 'emailaddress')

# Replace URLs with 'webaddress'

df_train['comment_text'] =
df_train['comment_text'].str.replace(r'^http:\/\/[a-zA-Z0-9\-\
\.] + \. [a-zA-Z]{2,3} (\S*)?$', 'webaddress')

# Replace money symbols with 'moneysymb' (£ can be typed
with ALT key + 156)
```

```

df_train['comment_text'] =
df_train['comment_text'].str.replace(r'£|\$', 'dollers')

# Replace 10 digit phone numbers (formats include
# paranthesis, spaces, no spaces, dashes) with 'phonenumber'

df_train['comment_text'] =
df_train['comment_text'].str.replace(r'^\d{3}\d{3}\d{4}[\s-]?[\d]{3}[\s-]?[\d]{4}$', 'phonenumber')

# Replace numbers with 'number'

df_train['comment_text'] =
df_train['comment_text'].str.replace(r'\d+(\.\d+)?',
'number')

# Remove punctuation

df_train['comment_text'] =
df_train['comment_text'].str.replace(r'[\w\d\s]', ' ')

# Replace whitespace between terms with a single space

df_train['comment_text'] =
df_train['comment_text'].str.replace(r'\s+', ' ')

# Remove leading and trailing whitespace

df_train['comment_text'] =
df_train['comment_text'].str.replace(r'^\s+|\s+?$', '')

```

2. Preparation for removal of punctuation marks & Updating the list of stop words:

```

df_train['comment_text'] =
df_train['comment_text'].apply(lambda x: ' '.join(
    term for term in x.split() if term not in
    string.punctuation))

stop_words = set(stopwords.words('english') + ['u', 'ü',
'ur', '4', '2', 'im', 'dont', 'doin', 'ure'])
df_train['comment_text'] =
df_train['comment_text'].apply(lambda x: ' '.join(
    term for term in x.split() if term not in stop_words))

lem=WordNetLemmatizer()
df_train['comment_text'] =

```

```
df_train['comment_text'].apply(lambda x: ' '.join(
    lem.lemmatize(t) for t in x.split()))
```

3. Stemming and Lemmatizing :

I imported the string library comprising all punctuation characters and appended the numeric digits to it, as those were required to be removed too.

Stemming is the process of converting inflected/derived words to their word stem or the root form. Basically, a large number of similar origin words are converted to the same word. E.g. words like “stems”, “stemmer”, “stemming”, “stemmed” are based on “stem”. This helps in achieving the training process with a better accuracy.

Lemmatizing is the process of grouping together the inflected forms of a word so they can be analyzed as a single item. This is quite similar to stemming in its working but not exactly same. Lemmatizing depends on correctly identifying the intended part of speech and meaning of a word in a sentence, as well as within the larger context surrounding that sentence, such as neighboring sentences or even an entire document. I used the **word-net library** in nltk for this purpose. Stemmer and Lemmatizer were imported from nltk.

4. Applying Tf-id Vectorizer :

Tf-id Vectorizer is used for converting a string of words into a matrix of words. Column headers have the words themselves and the cell values signify the frequency of occurrence of the word.

```
# Convert text into vectors using TF-IDF
tf_vec = TfidfVectorizer(max_features = 10000,
    stop_words='english')
X_features = tf_vec.fit_transform(df_train['comment_text'])
X = X_features
```

STEP 6: Model Building :


```

y = df_train['malignant']

# splitting into train and test

x_train,x_test,y_train,y_test=train_test_split(
X,y,random_state=56,test_size=.30)

#RandomForestClassifier
RF = RandomForestClassifier()

RF.fit(x_train, y_train)
y_pred_train = RF.predict(x_train)
print('Training accuracy is
{}'.format(accuracy_score(y_train, y_pred_train)))
y_pred_test = RF.predict(x_test)
print('Test accuracy is
{}'.format(accuracy_score(y_test,y_pred_test)))
print(confusion_matrix(y_test,y_pred_test))
print(classification_report(y_test,y_pred_test))

Training accuracy is 0.9988451105202374
Test accuracy is 0.9555899064171123
[[42695   552]
 [ 1574  3051]]

```

	precision	recall	f1-score	support
0	0.96	0.99	0.98	43247
1	0.85	0.66	0.74	4625
accuracy			0.96	47872
macro avg	0.91	0.82	0.86	47872
weighted avg	0.95	0.96	0.95	47872

```

# DecisionTreeClassifier
DT = DecisionTreeClassifier()

DT.fit(x_train, y_train)
y_pred_train = DT.predict(x_train)
print('Training accuracy is
{}'.format(accuracy_score(y_train, y_pred_train)))
y_pred_test = DT.predict(x_test)
print('Test accuracy is
{}'.format(accuracy_score(y_test,y_pred_test)))
print(confusion_matrix(y_test,y_pred_test))
print(classification_report(y_test,y_pred_test))

Training accuracy is 0.9988809210467416
Test accuracy is 0.9396933489304813

```

```

[[41867 1380]
 [ 1507 3118]]

```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	43247
1	0.69	0.67	0.68	4625
accuracy			0.94	47872
macro avg	0.83	0.82	0.83	47872
weighted avg	0.94	0.94	0.94	47872

```

# MultinomialNBClassifier
MN = MultinomialNB()

MN.fit(x_train, y_train)
y_pred_train = MN.predict(x_train)
print('Training accuracy is
{}'.format(accuracy_score(y_train, y_pred_train)))
y_pred_test = MN.predict(x_test)
print('Test accuracy is
{}'.format(accuracy_score(y_test, y_pred_test)))
print(confusion_matrix(y_test, y_pred_test))
print(classification_report(y_test, y_pred_test))

Training accuracy is 0.9514588313234675
Test accuracy is 0.9488218582887701
[[43079 168]
 [ 2282 2343]]

```

	precision	recall	f1-score	support
0	0.95	1.00	0.97	43247
1	0.93	0.51	0.66	4625
accuracy			0.95	47872
macro avg	0.94	0.75	0.81	47872
weighted avg	0.95	0.95	0.94	47872

```

# GradientBoostingClassifier
GBC = GradientBoostingClassifier()

GBC.fit(x_train, y_train)
y_pred_train = GBC.predict(x_train)
print('Training accuracy is
{}'.format(accuracy_score(y_train, y_pred_train)))

```

```

y_pred_test = GBC.predict(x_test)
print('Test accuracy is
{}'.format(accuracy_score(y_test,y_pred_test)))
print(confusion_matrix(y_test,y_pred_test))
print(classification_report(y_test,y_pred_test))

Training accuracy is 0.9988809210467416
Test accuracy is 0.9396933489304813
[[41867  1380]
 [ 1507   3118]]

```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	43247
1	0.69	0.67	0.68	4625
accuracy			0.94	47872
macro avg	0.83	0.82	0.83	47872
weighted avg	0.94	0.94	0.94	47872


Plotting the graph which tells us about the area under curve , more the area under curve more will be the better prediction

model is performing good :

```

fpr,tpr,thresholds=roc_curve(y_test,y_pred_test)
roc_auc=auc(fpr,tpr)
plt.plot([0,1],[1,0], 'k--')
plt.plot(fpr,tpr,label = 'RF Classifier')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.title('RF CLASSIFIER')
plt.show()

```



as per found above models scores random forest is giving better accuracy hence considering as final model and saving it using joblib library.

Predicting the test data:

```
df_test =tf_vec.fit_transform(df_test['comment_text'])
df_test

prediction=RF.predict(df_test)
prediction
```

STEP 7: Saving Model :

```
import joblib

joblib.dump(RF,"malignant_commant.pkl")

['malignant_commant.pkl']
```

STEP 7: Conclusion :

To conclude, In this project we have performed various techniques like Data analysis which has proven way for any organizations and

enterprises to gain the information they need to make better decisions, serve their customers, and increase productivity and revenue. The benefits of data analysis are almost too numerous to count, and some of the most rewarding benefits include getting the right information for business, getting more value out of it. And it creates more effective marketing campaigns, gaining a better understanding of customers, and so on. In addition, Stemming is the process of converting inflected/derived words to their word stem or the root form. Basically, a large number of similar origin words are converted to the same word. E.g. words like “stems”, “stemmer”, “stemming”, “stemmed” are based on “stem”. This helps in achieving the training process with a better accuracy.

Lemmatising is the process of grouping together the inflected forms of a word so they can be analyzed as a single item. This is quite similar to stemming in its working but not exactly same. Lemmatising depends on correctly identifying the intended part of speech and meaning of a word in a sentence, as well as within the larger context surrounding that sentence, such as neighboring sentences or even an entire document. I used the word-net library in nltk for this purpose. Stemmer and Lemmatizer were imported from nltk.

Tfidf Vectorizer is used for converting a string of words into a matrix of words. Column headers have the words themselves and the cell values signify the frequency of occurrence of the word. Finally I reached the core part of the project, where I could start building the classifier. I had use 4 major algorithm . Where, our random forest model giving the good accuracy hence selected it and saved it using joblib library.